



# SDA-CAPSTONE – Blog App Deployment

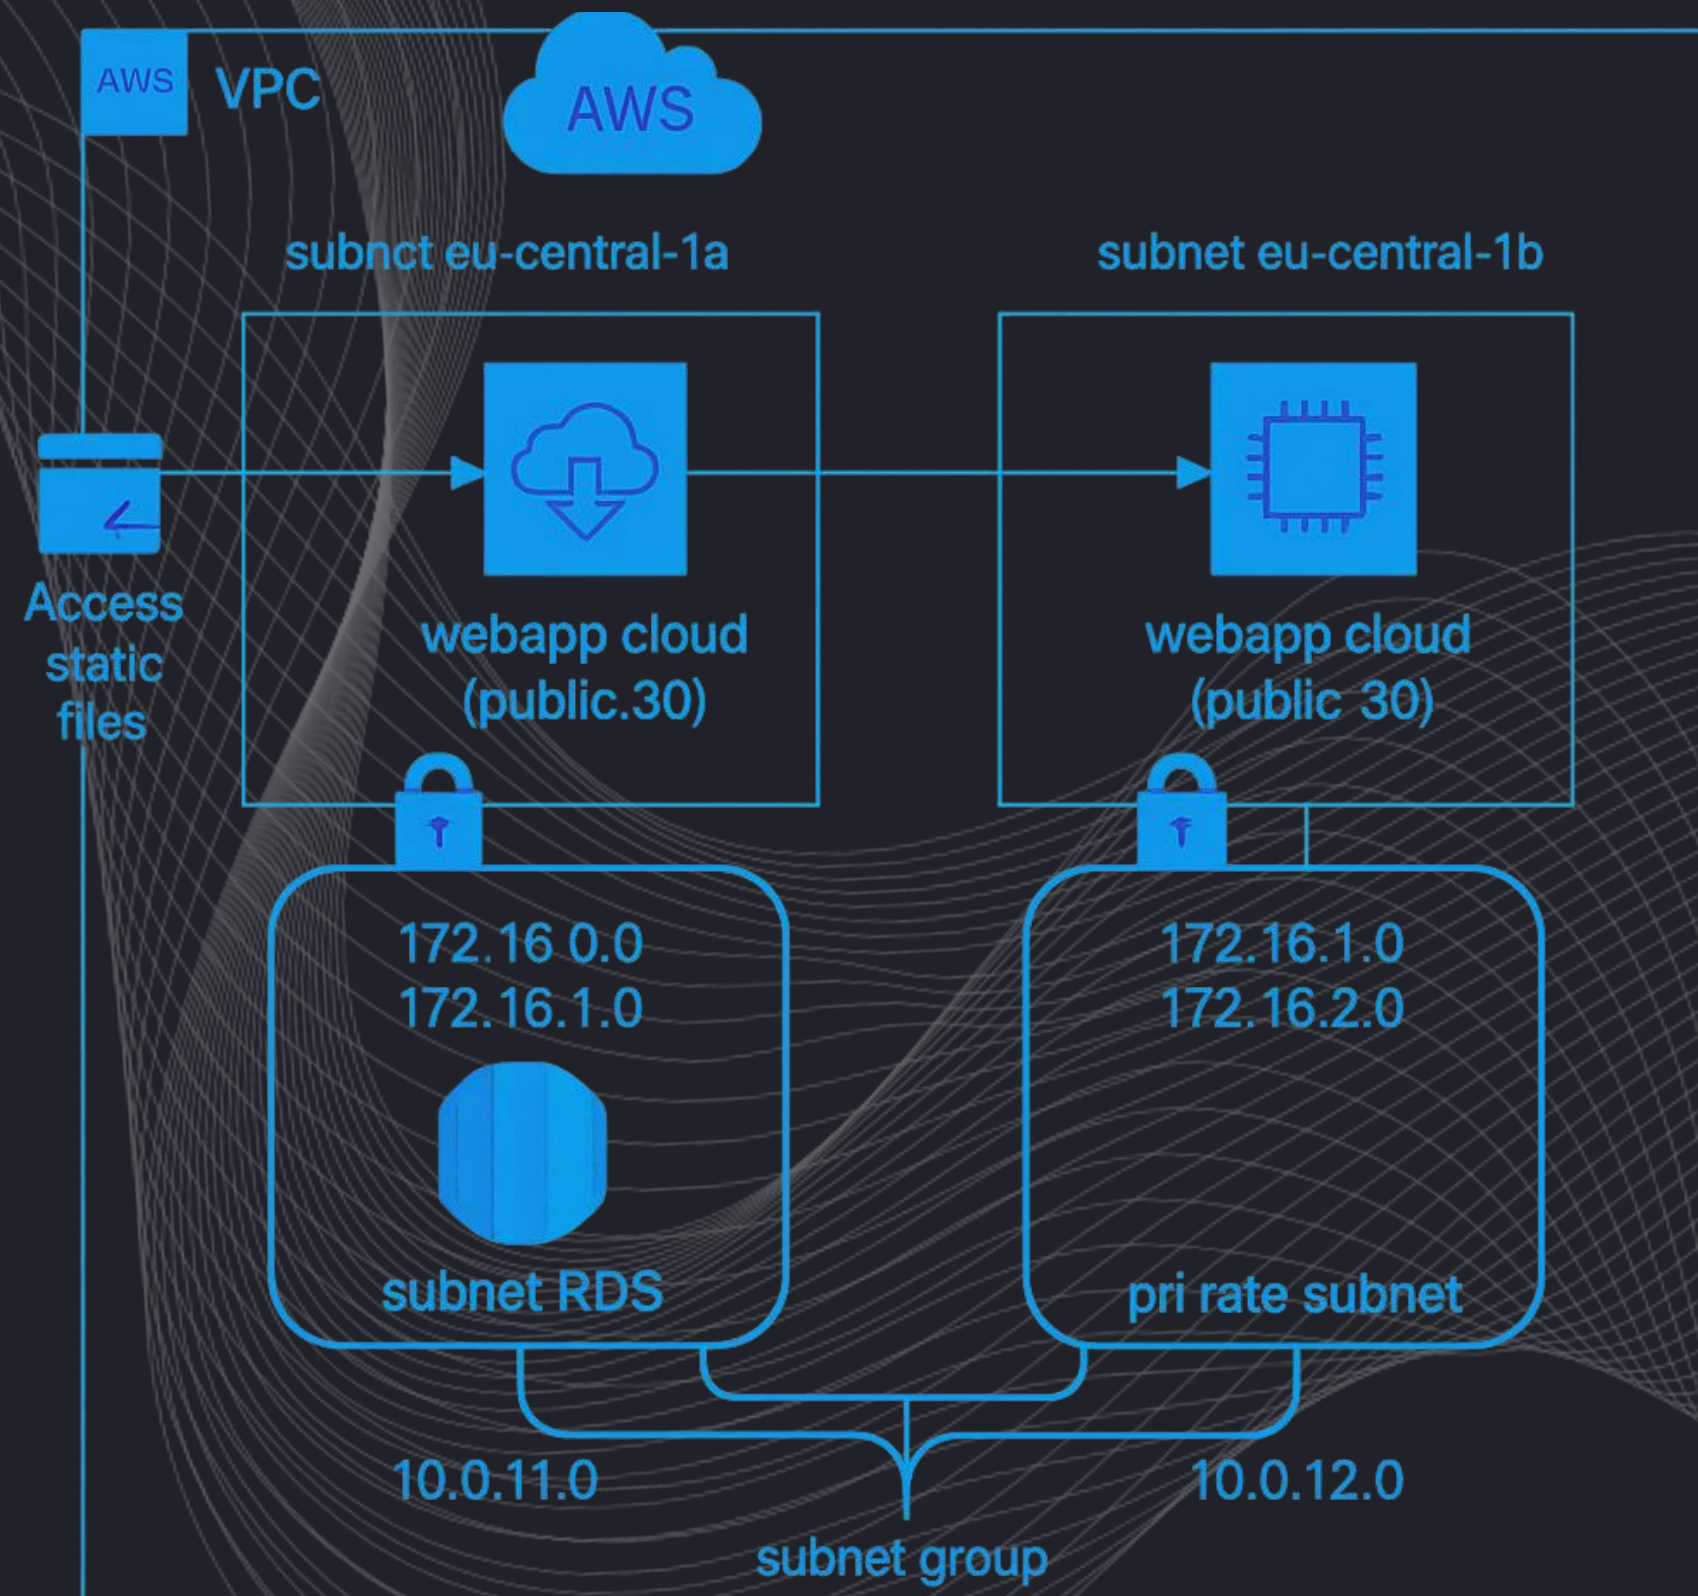
## Team 1

Ibrahim-SDA2002   Manal-SDA2020   Shadan-SDA2005   Manal-SDA2006   Bushra-SDA2015   Rawan-SDA2022  
Renad-SDA2038   Khalid-SDA2024   Lina-SDA2027   Mafaz-SDA2029   Nawal-SDA2039   Huda-SDA2001



# Architecture Overview

SDA



## Overview of Architecture Components

Custom VPC with public and private subnets distributed across two Availability Zones (eu-central-1a & eu-central-1b).




# Networking Architecture

SDA

What to do :

- Create VPC: 10.90.0.0/16
- 2 Public & 2 Private subnets across 2 AZs
- Enable DNS hostnames

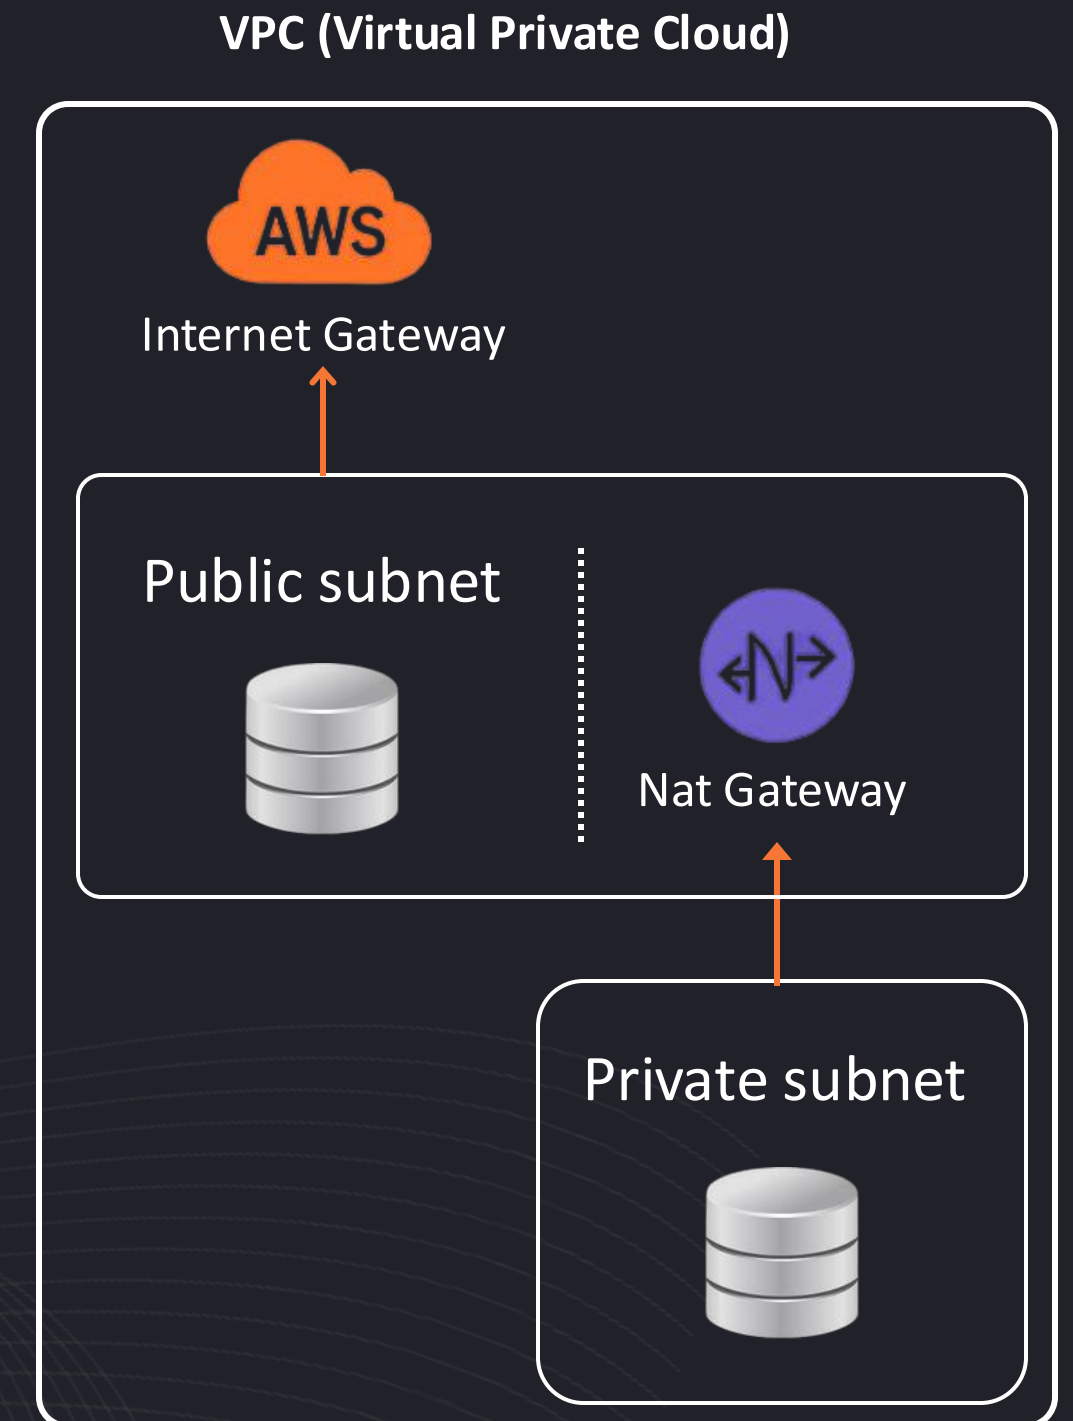


# Internet & NAT Gateway Setup

SDA

- **Public subnets** are connected to the Internet Gateway (IGW) to allow direct internet access for resources like web servers.
- **Private subnets** route outbound traffic through a NAT Gateway, which enables instances in private subnets to access the internet securely without exposing them to inbound traffic.
- **Route Tables** are configured to direct traffic properly:
  - Public subnets → route via IGW for inbound and outbound internet traffic.
  - Private subnets → route via NAT Gateway for outbound internet access only.

This setup improves security by isolating private resources from direct internet exposure while maintaining their ability to update and communicate externally.





# Security Groups

SDA



**ALB SG:** HTTP/HTTPS from anywhere

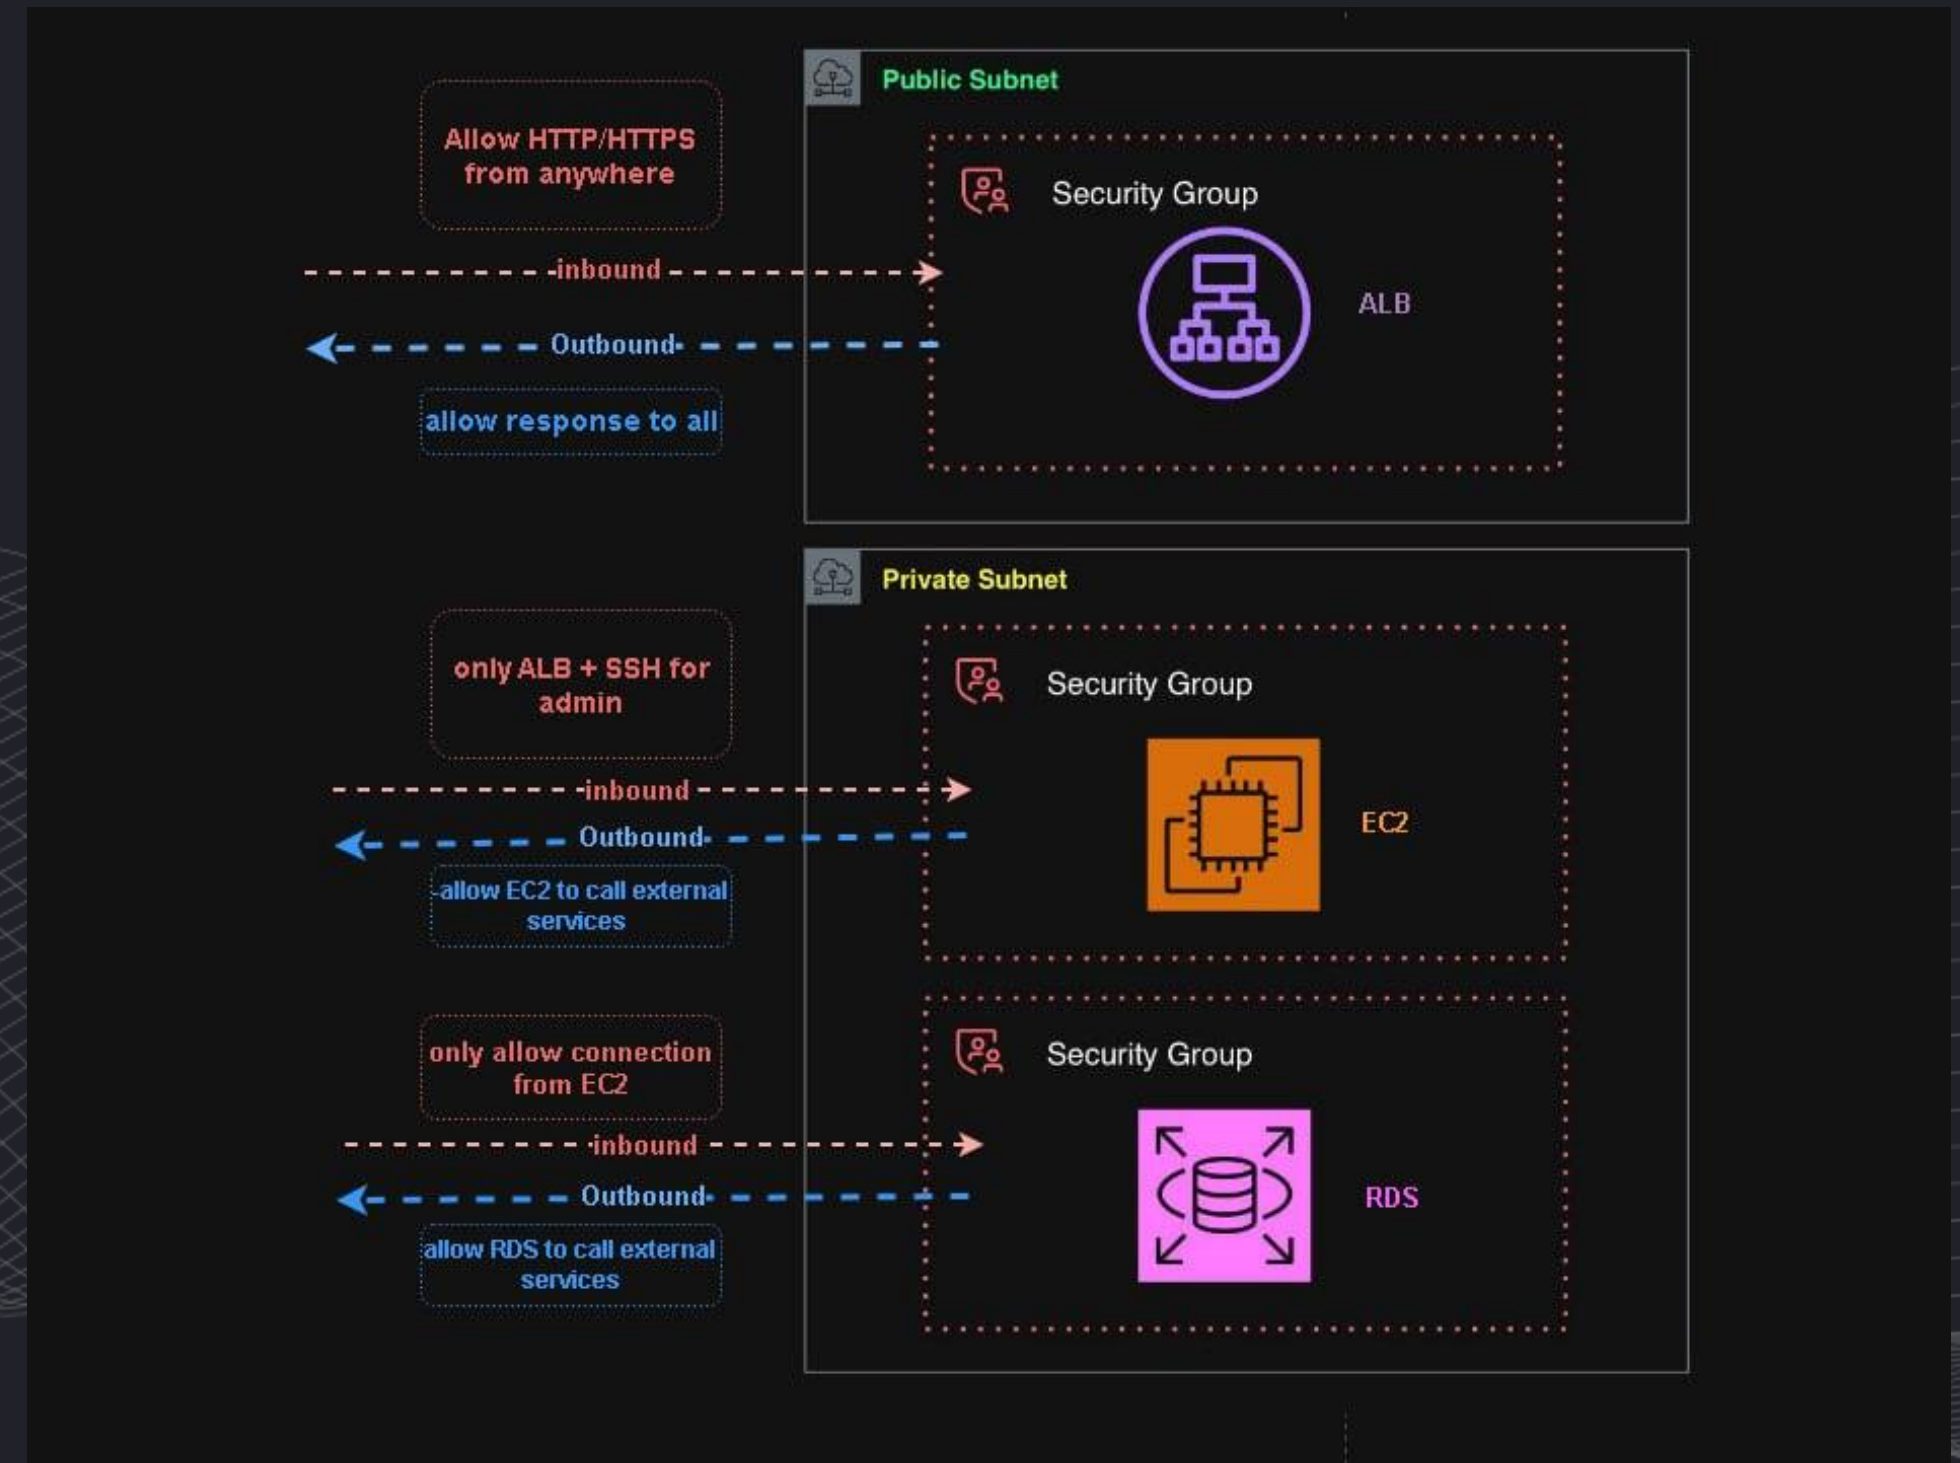
Allows HTTP/HTTPS from anywhere for public access

**EC2 SG:** Only ALB + SSH for admin

Allows traffic only from the ALB, plus SSH for admin access.

**RDS SG:** Only EC2

Allows only EC2 instances to connect to the database.





# GitHub Integration

---

SDA



Private GitHub repo for source code



Token-based clone using HTTPS



Automate file changes with `github_repository_file` in Terraform

Store sensitive values:

/capstone/username

/capstone/password

/capstone/token

These are stored securely in AWS Systems Manager Parameter Store as SecureStrings.

## Used by EC2 & Terraform

### EC2 Instances:


Use `aws ssm get-parameters` to fetch secrets at runtime for environment variables or config files.

### Terraform:

Fetches SSM parameters using the `aws_ssm_parameter` data source to inject secrets into infrastructure as code (e.g., for user data or app setup).



AWS  
SSM





# RDS Configuration

---

SDA



Private subnet  
group  
(2 A2s)

MySQL &.x.  
db.t4g.micro

Not publicly  
accessible

Use SSM for  
DB credentials



# S3 Bucket for Media

SDA

Bucket : name it as you like

ACL enabled, public access allowed

Stores uploaded images/videos  
from the app

Up and ready





## EC2 Test Instance :

This EC2 instance was created just for testing purposes.  
We launched it inside a public subnet so we could access it easily during setup.

We attached an IAM role that gives it access to:

- S3(to store and retrieve media files),
- SSM(to connect securely without SSH),
- and RDS(our database).



Amazon  
EC2

After launching the instance, we manually ran the userdata.sh script to:

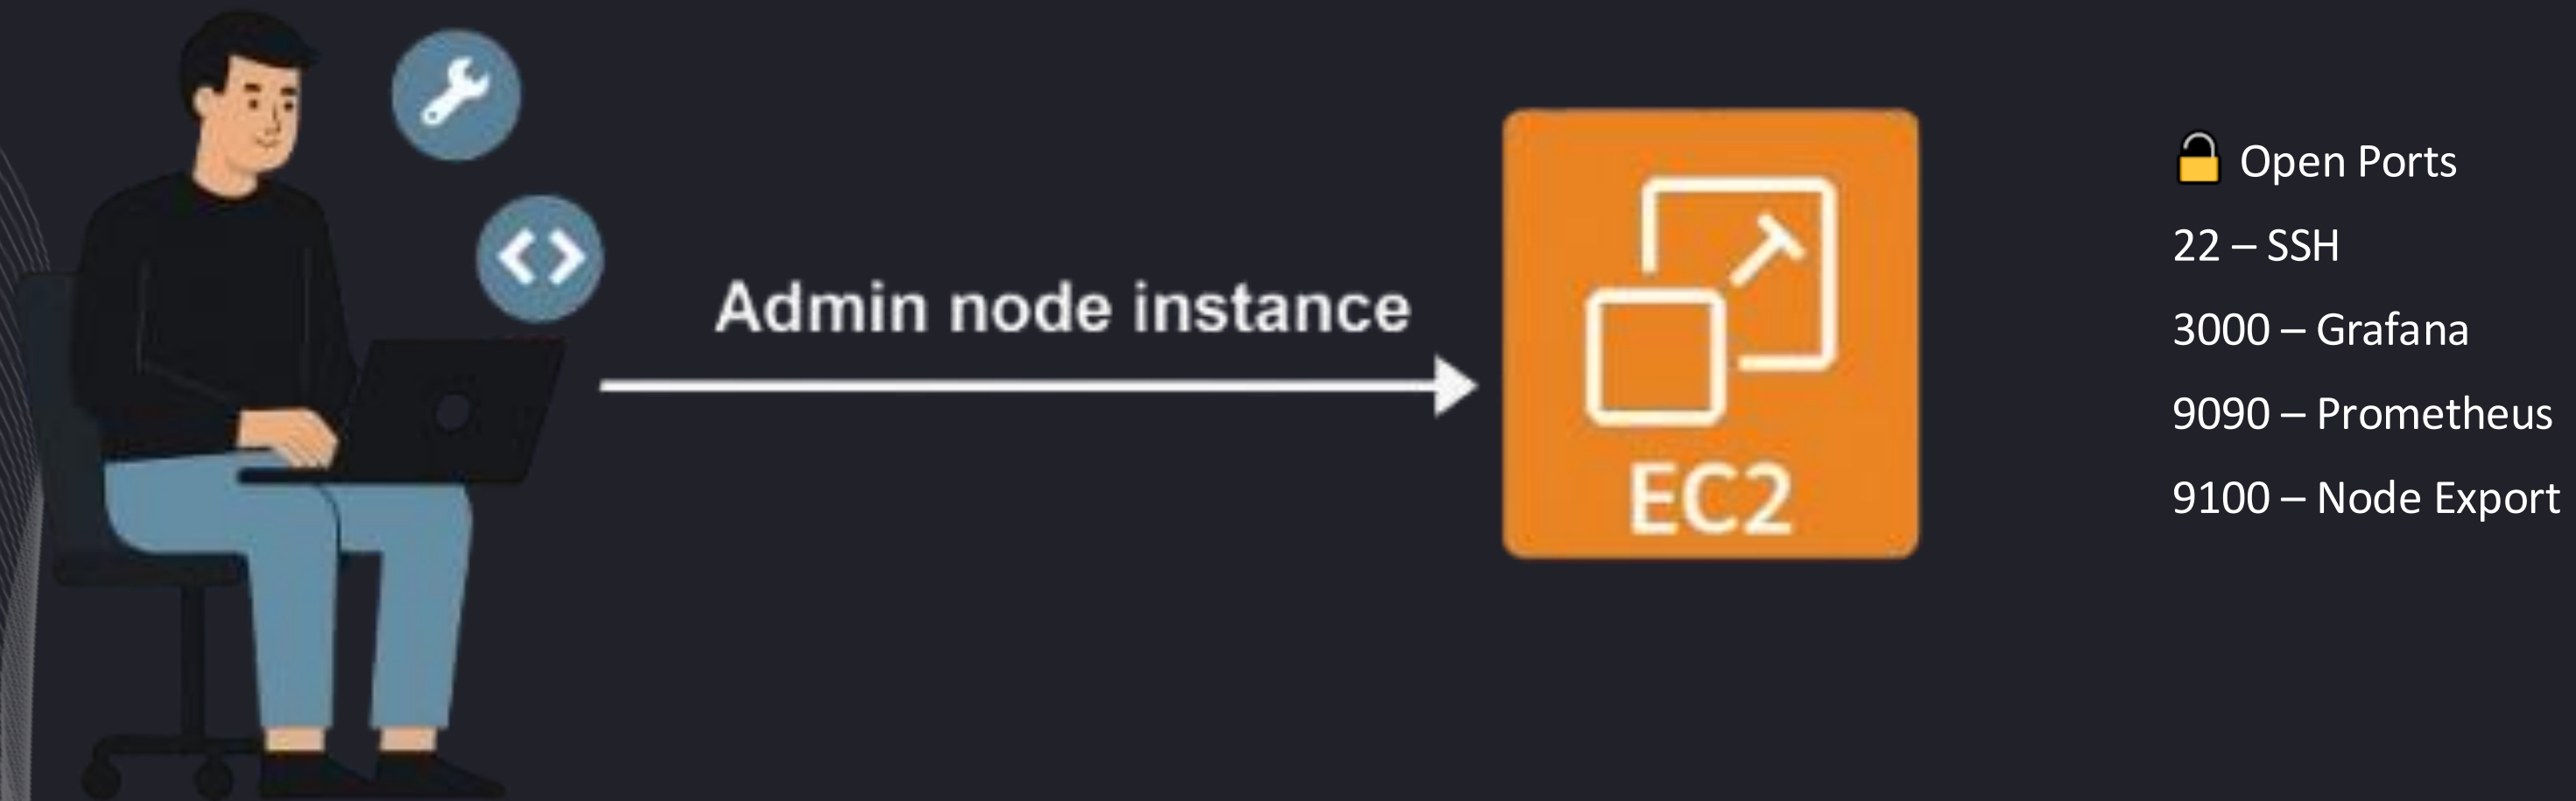
- test the Django installation,
- check if environment variables work properly,
- and make sure our server setup is complete.

This helped us confirm that everything – like the app, database, and storage – is connected and working  
before moving to productio



# Admin node setup

SDA



🔓 Open Ports  
22 – SSH  
3000 – Grafana  
9090 – Prometheus  
9100 – Node Export

## ◆ Central EC2 Instance (Amazon Linux)

Used for:

- JumpBox – secure entry point for SSH access
- Ansible Control Node – configuration management and automation
- Monitoring Stack – runs Prometheus and Grafana for system metrics and dashboards



## Resources

**3**

GitHub automation  
for settings.py

**2**

IAM Roles, Security  
Groups, RDS

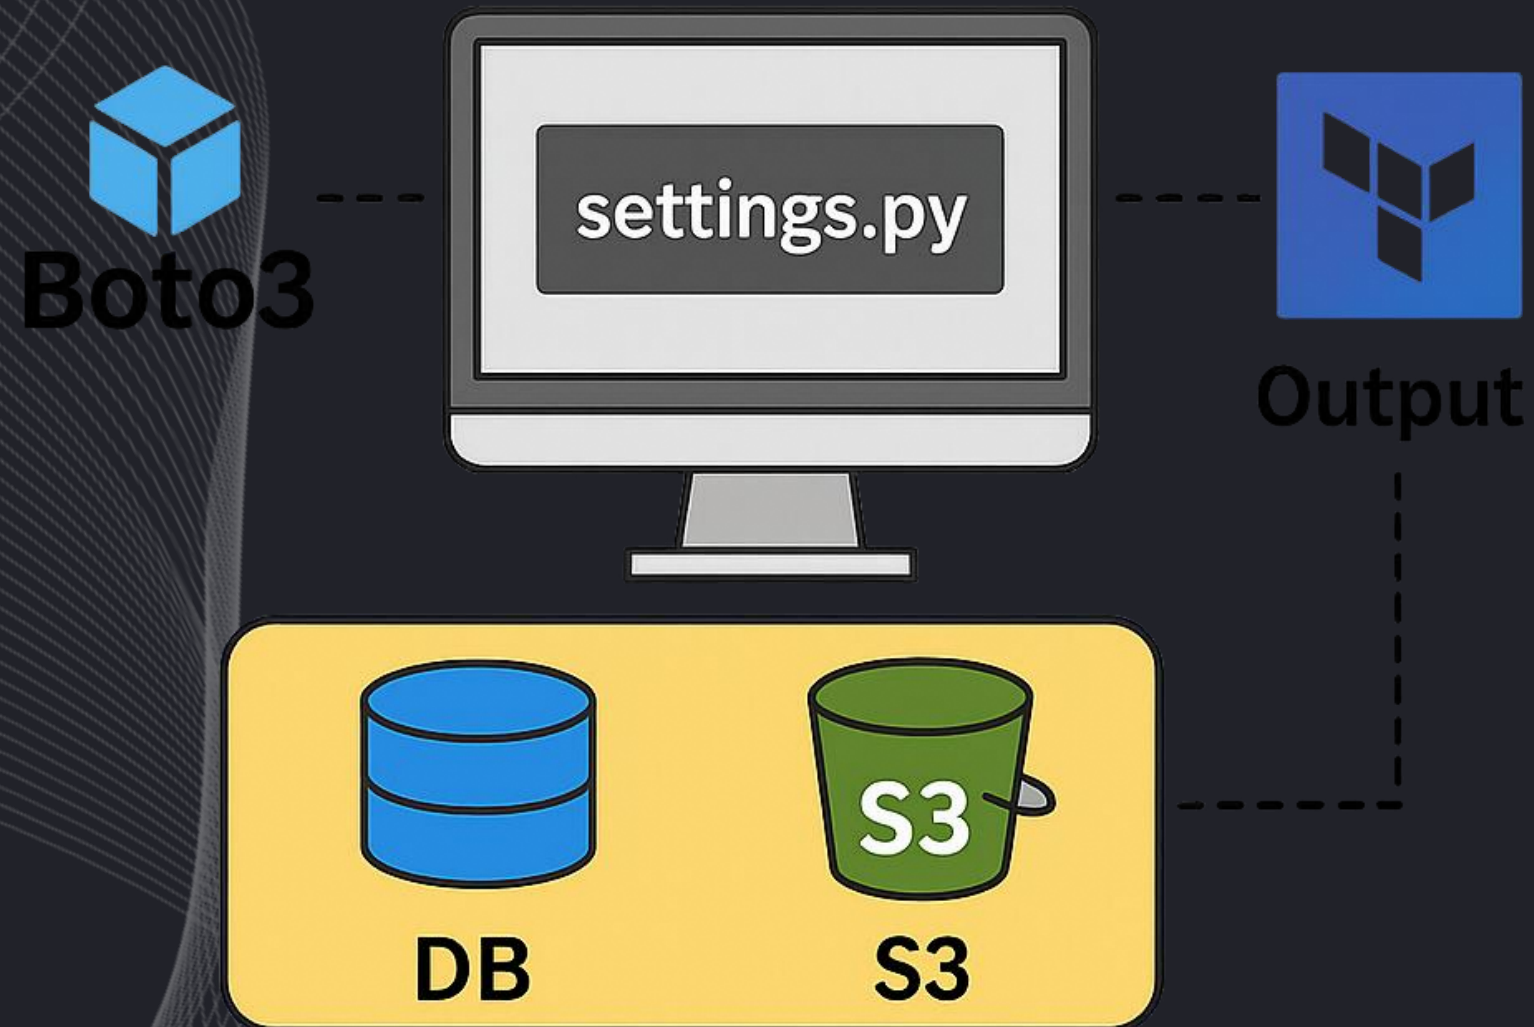
**1**

ALB, Target Group, Launch  
Template, ASG



# Django settings.py Configuration

SDA




## 🔑 Key Points for settings.py Configuration

Boto3 is used to fetch secrets (DB credentials, tokens) securely from AWS SSM Parameter Store

Terraform output dynamically injects the DB endpoint into the Django config  
S3 and DB credentials are loaded at runtime via environment variables, keeping settings.py clean and secure

## Ansible & Dynamic Inventory



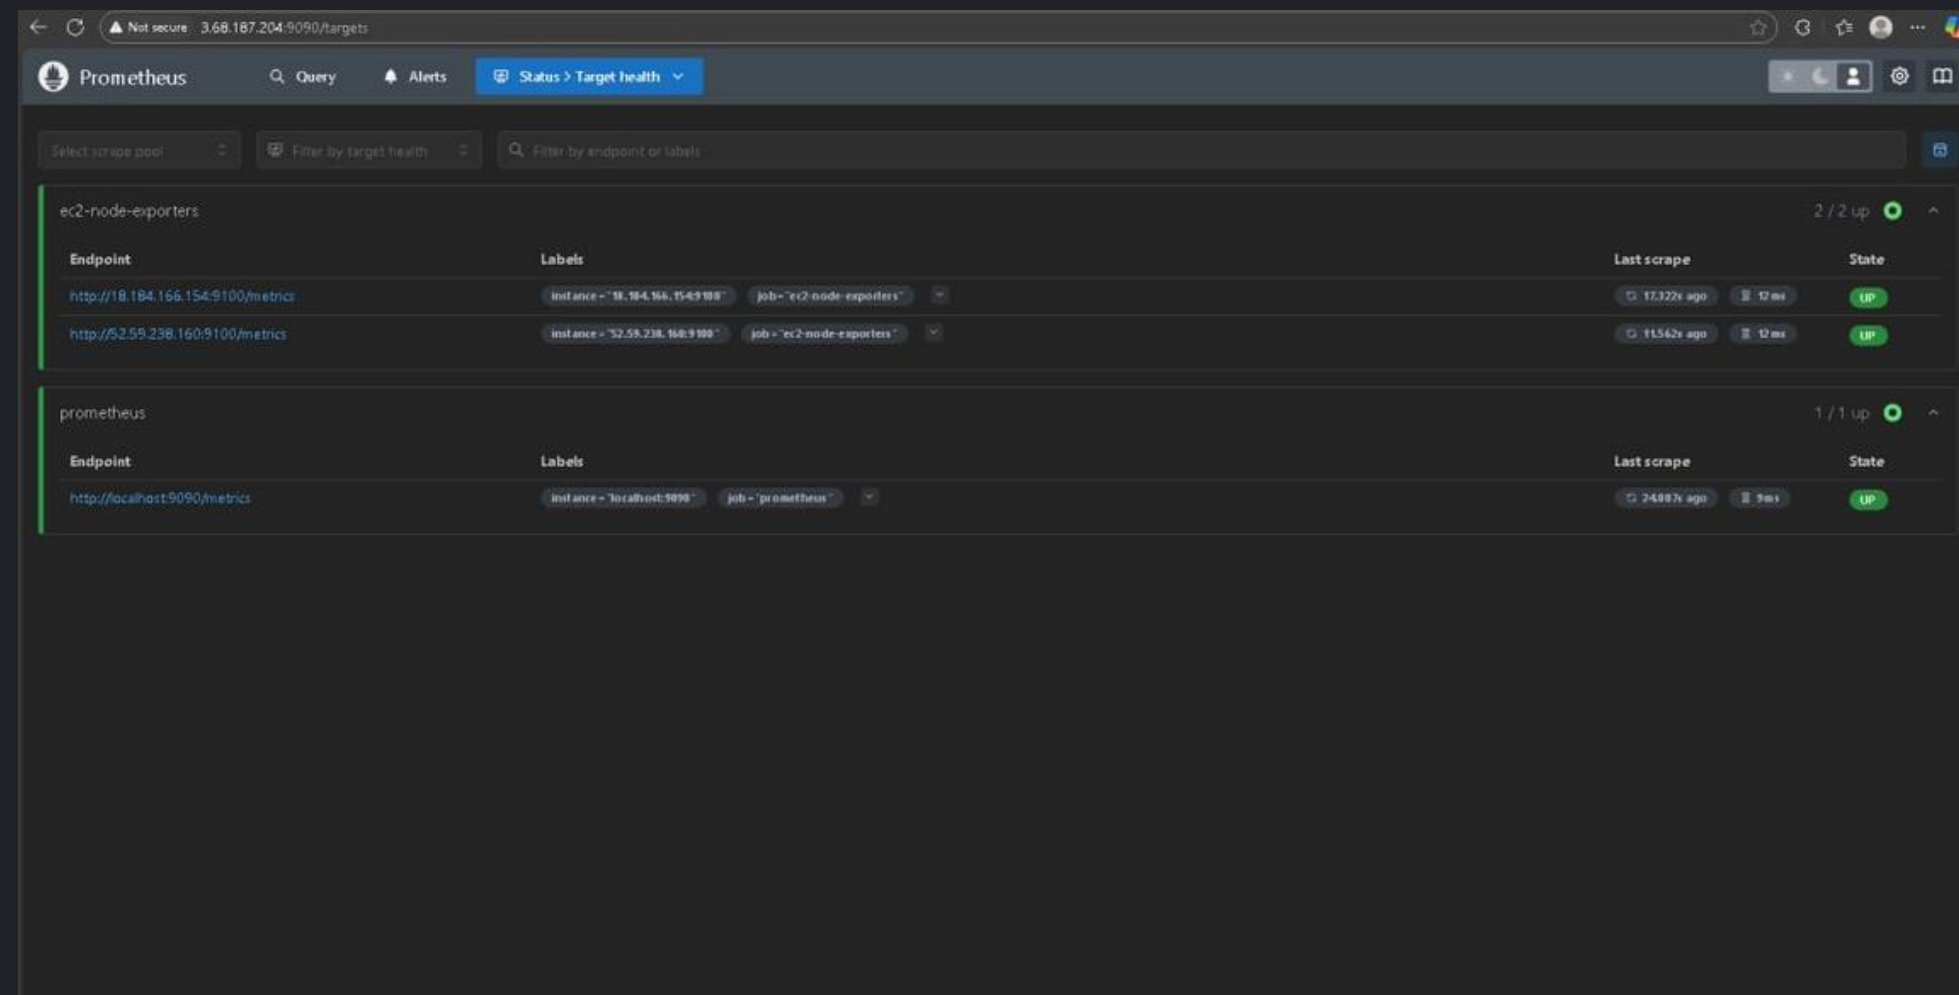
- Use AWS EC2 dynamic inventory plugin for real-time host discovery
- Filter instances by tags:
- Project=Capstone, App=BlogPage
- Set up ansible.cfg and provide the .pem key for secure access



# Monitoring with Prometheus

SDA

- Deploy Node Exporter on EC2 instances using Ansible
- Enable EC2 service discovery in Prometheus (ec2\_sd\_configs)
- Monitor key metrics: CPU usage, memory, disk space, and uptime



# Grafana dashboard setup:

---

SDA

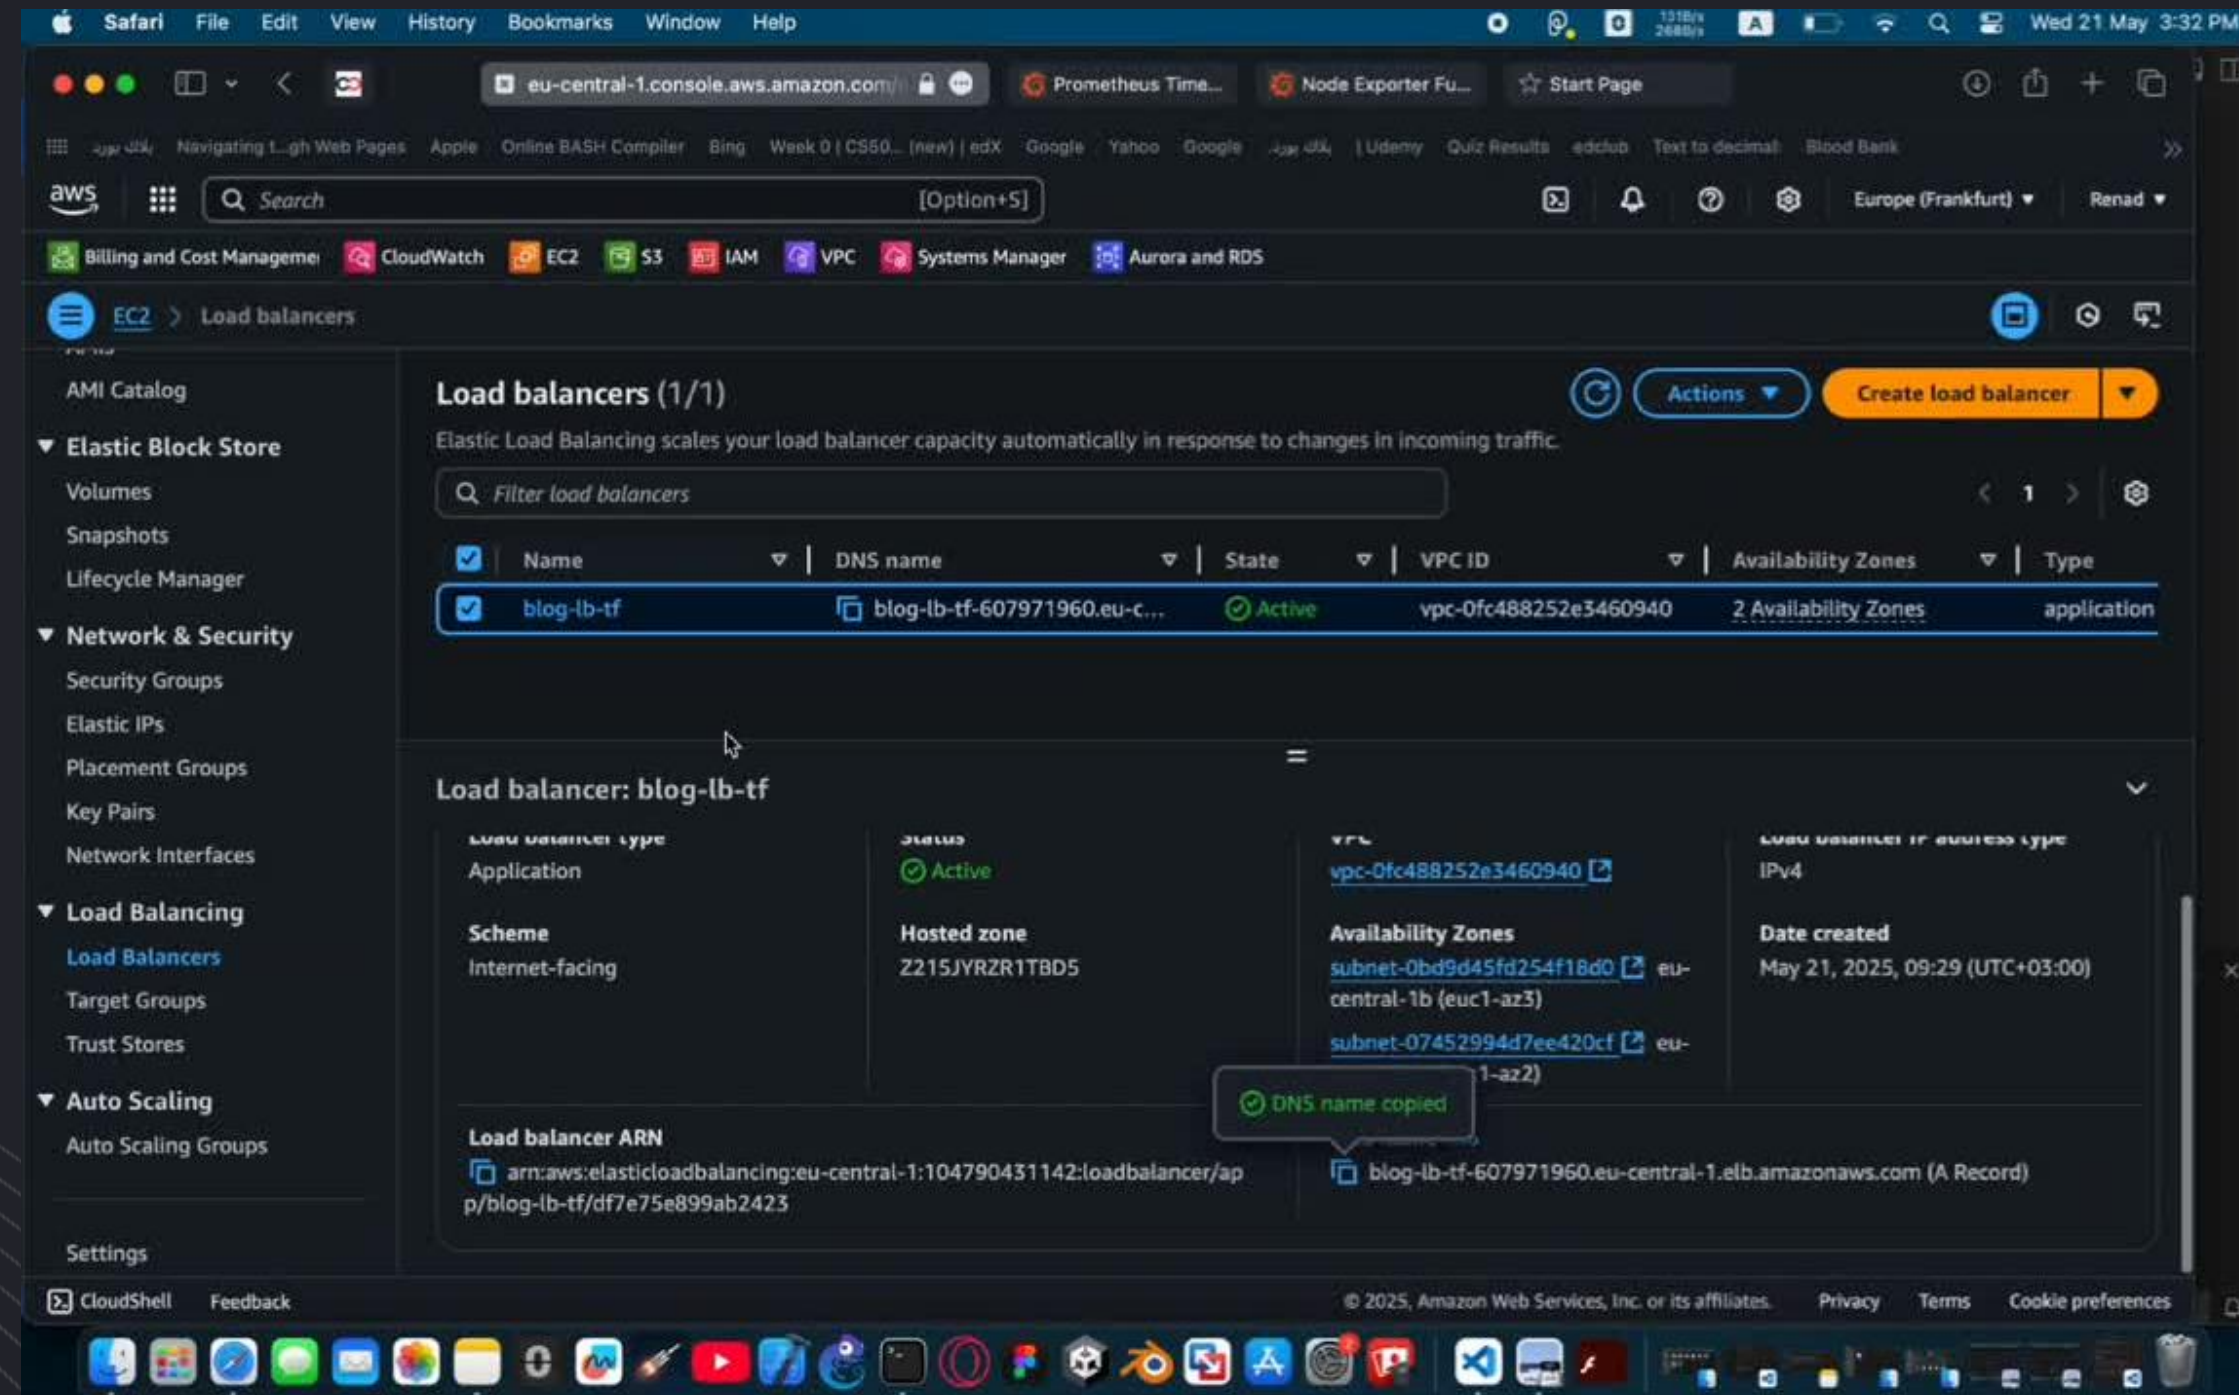
**Log in to Grafana with default  
credentials: admin / admin**

**Add Prometheus as a data source  
via server URL**  
`http://localhost:9090`

**Import Dashboard ID 1860 – Node  
Exporter Full for system metrics**



## Time for a live demonstration !!!





## challenges

- **The initial AMI image was incompatible with our region, so we had to search for and configure a suitable one manually.**
- **Encountered several IAM permission errors that required time-consuming debugging and policy adjustments.**
- **Some Terraform modules were outdated or misconfigured, leading to deployment failures and resource rollbacks.**
- **Managing the SSH access for multiple teammates required careful coordination and key sharing.**
- **File permission issues on the EC2 instance caused problems during Ansible deployment.**
- **Environment variables were not being picked up properly by the backend service initially, delaying testing.**





# Thank You

GitHub Repositor:

[https://github.com/NawalSuf/infrastructure-  
capstone-  
team1/blob/main/README.md](https://github.com/NawalSuf/infrastructure-capstone-team1/blob/main/README.md)